



Lab Exercise: Networking Containers

Preamble

This lab exercise demonstrates how it is possible to control a container's networking capabilities, including the forwarding of a service port to the host, and how a container's service can be consumed from the Docker host, a sibling container, or from the container itself. This will be performed on a bridge network.

It also demonstrates how to establish a multi-host overlay network, and illustrates inter-host container service consumption. The lab exercise requires multiple Docker hosts, which can be achieved with Oracle Virtualbox VMs.

Bridge Network

Step 1 – create a new user-defined bridge network

List the currently available Docker networks registered with your local Docker daemon.

```
$ docker network ls
```

You should see output similar to the following:

NETWORK ID	NAME	DRIVER
a35b69b5ed3e	none	null
76a226c6cad7	host	host
913227fccf96	bridge	bridge

Using the docker network client command, create a new bridge network called 'my_bridge', with a subnet of 172.20.0.0/16 and a gateway of 172.20.0.254.

```
$ docker network create --subnet 172.20.0.0/16 --gateway 172.20.0.254 \  
> my_bridge
```

Use the docker network ls sub-command to confirm it has been created, and then use the docker network inspect sub-command to check it has been configured correctly.

Step 2 – forward a container's port

Start a new container based on the Ubuntu library image, connected to the 'my_bridge' network, give it the name 'netcat', and specify a port to forward to the host:

```
$ docker run -it --rm --name netcat--net my_bridge -p :80 ubuntu:latest \
> /bin/bash
```

Check which host port the container's port 80 has been forwarded to using the docker inspect command:

```
$ docker inspect -f \
> '{{(index (index .NetworkSettings.Ports "80/tcp") 0).HostPort}}' \
> netcat
32769
```

At the bash command shell inside the container, start the netcat utility listening on port 80:

```
$ nc -k -l 80
```

Start a new bash command shell on the host, and invoke the netcat utility, using the port number retrieved with the docker inspect command:

```
$ nc $(hostname) 32769
```

Type some text in the bash command shell on the host, and verify that it echoed within the container (remember to hit the <ENTER> key).

Terminate the netcat utility running on the host by typing CTL-C, but leave netcat running in the container.

Step 3 – connect from another container

Start a new container based on the Ubuntu library image, with a bash command shell as the container's process:

```
$ docker run -it --rm ubuntu:latest /bin/bash
```

From within the container, invoke the netcat utility, using the same port number you used in the previous step, using the hostname associated with your host:

```
$ nc <hostname> 32769
```

Type some text in the bash command shell, and verify that it is echoed within the container whose port has been forwarded to the host. Notice that the container was able to lookup the IP address of the host via the embedded DNS server provided by the Docker daemon.

Terminate the netcat utility running in the container by typing CTL-C, and exit the container (not the container with the port forwarded).

Step 4 – connect from the same container

Finally, we're going to use the netcat utility to make the container consume its own service.

From a bash command shell on the host, execute an interactive bash command shell in our container with the port forwarded, using the docker exec command:

```
$ docker exec -it netcat /bin/bash
```

In the new command shell, use the netcat utility to connect to the socket on the host (substituting <hostname> for the actual hostname, and <host port no.> for the actual port number):

```
$ nc <hostname> <host port no.>
```

Type a message and verify that it is echoed within the container that we have been monitoring in the previous steps.

Terminate the netcat utility running in both containers by typing CTL-C, and exit both containers. Use the docker network rm sub-command to remove the 'my_bridge' network from the list of networks available to connect to.

Overlay Network

To demonstrate the use of overlay networks in Docker, we will use a variation of Dj Walker-Morgan's [goredchat](#) application, a simple chat application that uses the Redis database engine to register chat users, and for routing chat messages. We'll create a Redis instance, and two client sessions using goiredchat (all running in containers), but on different Docker hosts and connected via an overlay network.

Step 1 – establish a key/value store

The first thing we need to do is establish the key/value store that overlay networking requires. We will use Hashicorp's Consul key/value store, and the easiest way to do this is to run it in a container on a dedicated VM. To simplify things, we'll use Docker Machine to create the VM. If docker-machine is not already installed locally, then follow [these instructions](#).

Create the VM:

```
$ docker-machine create -d virtualbox kv-store
```

Point your Docker client at the Docker host running on the kv-store VM:

```
$ eval $(docker-machine env kv-store)
```

Check this has taken effect by issuing the following command, which should return the result shown:

```
$ docker info | grep Name
Name: kv-store
```

The next step is to start a container running Consul, and we'll use the `progrium/consul` image from the Docker Hub. The image needs some ports forwarded to its VM host, and can be started with the following command:

```
$ docker run -it -d --restart unless-stopped -p 8400:8400 -p 8500:8500 \
> -p 8600:53/udp -h consul progrium/consul -server -bootstrap -ui-dir /ui
```

Check that Consul is up and running by determining the IP address of the kv-store VM, and then entering the following URL into a web browser; `<ip>:8500/ui`:

```
$ docker-machine ip kv-store
192.168.99.100
```

If all is well, you will see the user-interface for the Consul application.

Step 2 – create three Docker hosts additional Docker hosts

Next, we'll create three different VMs, each running a Docker daemon, which we'll use to host containers that will be connected to the overlay network. Each of the Docker daemons running on these machine needs to be made aware of the KV store, and of each other.

```
$ for i in {1..3}; do \
> docker-machine create -d virtualbox --engine-opt \
> "cluster-store consul://$(docker-machine ip kv-store):8500" \
> "cluster-advertise eth1:2376" \
> host0${i}; done
```

Check that all the VMs are running as expected:

```
$ docker-machine ls -f \
> "table {{.Name}} \t{{.State}}\t{{.URL}}\t{{.DockerVersion}}"
NAME      STATE      URL                                DOCKER
host01     Running    tcp://192.168.99.104:2376         v1.10.2
host02     Running    tcp://192.168.99.105:2376         v1.10.2
host03     Running    tcp://192.168.99.106:2376         v1.10.2
kv-store   Running    tcp://192.168.99.100:2376         v1.10.2
```

Step 2 – create an overlay network

We now need to run some Docker CLI commands on each of the Docker hosts we have

created. There numerous ways of doing this;

1. Establish an ssh session on the VM in question, using the `docker-machine ssh` command
2. Point our local Docker client at the Docker host in question, using the `docker-machine env` command
3. Run one-time commands against the particular Docker host using the `docker-machine config` command

The overlay network can be created using any of the Docker hosts:

```
$ docker $(docker-machine config host01) network create -d overlay \
> my_overlay
0223fc182bd363b95fce819bd33640c7afff8f5b306f0a93a7093b9a254cd9cc
```

Check that the overlay network 'my_overlay' can be seen from each of the Docker hosts:

```
$ docker $(docker-machine config host01) network ls -f name=my_overlay
NETWORK ID          NAME                DRIVER
0223fc182bd3        my_overlay         overlay
```

Step 3 – create a container running the Redis KV database engine

First, we need to start a Redis server running in a container on Docker host host01, using the library image found on the Docker Hub registry.

```
$ docker $(docker-machine config host01) run -d --restart unless-stopped \
> --net-alias redis_svr --net my_overlay redis:alpine redis-server \
> --appendonly yes
```

The library Redis image will be pulled from the Docker Hub registry, and the Docker CLI will return the container ID, e.g.

```
4d96bb2bd436bb6a19d5408de026904adc91892db40cfa25a8d71bae0e73bb29
```

The Redis library image exposes port 6379, which will be accessible to containers running on the 'my_overlay' network. To test whether the `redis_svr` container is listening for connections, we can run the Redis client in an ephemeral container, also on host01:

```
$ docker $(docker-machine config host01) run --rm --net my_overlay \
> redis:alpine redis-cli -h redis_svr ping
PONG
```

Notice that Docker's embedded DNS server resolves the `redis_svr` name, which is an alias for the container, and the Redis server responds to the Redis client ping, with a PONG.

Step 4 – create a container running the goredchat application on host02

Now that we have established that a Redis server container is running on the 'my_overlay' network, we can attempt to consume its service from a container running on a different Docker host.

The goredchat image can be found on the Docker Hub registry, and can be run like a binary, with command line options. Run the following command to determine its usage:

```
$ docker $(docker-machine config host02) run --rm --net my_overlay \  
> nbrown/goredchat --help  
Usage: /goredchat [-r URL] username  
e.g. /goredchat -r redis://redis_svr:6379 antirez
```

If -r URL is not used, the REDIS_URL env must be set instead

Now run a container using the -r configuration option to address the Redis server:

```
$ docker $(docker-machine config host02) run -it --rm --net my_overlay \  
> nbrown/goredchat -r redis://redis_svr:6379 bill
```

```
Welcome to goredchat bill! Type /who to see who's online, /exit to exit.
```

Find out who's online by typing /who:

```
/who  
bill
```

The goredchat client has created a TCP socket connection from the container on host02 to the Redis server container on host01 via the 'my_overlay' network.

Step 5 – create a container running the goredchat application on host03

In another bash command shell, we can start another instance of goredchat, this time on host03.

```
$ docker $(docker-machine config host03) run -it --rm --net my_overlay \  
> nbrown/goredchat -r redis://redis_svr:6379 brian
```

```
Welcome to goredchat brian! Type /who to see who's online, /exit to exit.
```

Notice that in goredchat on host02, Bill gets a message that 'brian has joined'. Exchange some messages between Bill and Brian before exiting from the chat with /exit.

Stop and remove the Redis container, remove the 'my_overlay' network, stop and remove the Consul container, and then use docker-machine rm to remove the VMs.